

CS108: Advanced Database

- Database Programming

Lecture 02:

The Foundation Statements
of T-SQL

Overview

- The fundamental Transact-SQL (T-SQL) statements
 - SELECT
 - INSERT
 - UPDATE
 - DELETE

Fundamental T-SQL statements

- SELECT, INSERT, UPDATE, and DELETE are the bread and butter of T-SQL
- These statements make up the basis of T-SQL's Data Manipulation Language (DML)
- SQL also provides many operators and keywords that help refine the queries

Basic SELECT Statement

- The **SELECT** statement and the structures used are basic commands will perform with SQL Server
- The basic syntax rules for a **SELECT** statement:

```
SELECT [ALL|DISTINCT] [TOP (<expression>) [PERCENT] [WITH TIES]]  
      <column list>  
[FROM <source table(s)/view(s)>]  
[WHERE <restrictive condition>]  
[GROUP BY <column name or expression using a column in the SELECT list>]  
[HAVING <restrictive condition based on the GROUP BY results>]  
[ORDER BY <column list>]  
[[FOR XML {RAW|AUTO|EXPLICIT|PATH [(<element>)]}[, XMLDATA]  
                                     [, ELEMENTS][, BINARY base 64]]  
[;]
```

Elements of the SELECT Statement

- The clauses are logically processed in the following order:
 1. FROM - Queries the rows from ...
 2. WHERE - Filters only record where ...
 3. GROUP BY - Groups the record by ...
 4. HAVING - Filters only groups having ...
 5. SELECT - Selects (returns) for each group the ...
 6. ORDER BY - Orders (sorts) the rows in the output by ...

Basic Query

- A SELECT indicates that we are merely reading information, as opposed to modifying it.
- What we are selecting is identified by an expression or column list immediately following the SELECT.
- The FROM statement specifies the name of the table or tables from which we want to get our data.
- With these, we have enough to create a basic SELECT statement.

Example: Query Metadata

- Most modern relational databases store metadata in system tables
- In SQL Server, we can query metadata through system views or stored procedures
- INFORMATION_SCHEMA
 - An information schema view is a common method for obtaining metadata.
 - This schema is contained in each database. Each information schema view contains metadata for all data objects stored in that particular database.

Example: INFORMATION_SCHEMA

- Ways to obtain metadata
 - Standard SQL SELECT
 - System stored procedures
- Commonly used views
 - TABLES
 - COLUMNS
 - TABLE_CONSTRAINTS
 - VIEWS
 - KEY_COLUMN_USAGE

	TABLE_CATALOG	TABLE_SCHEMA	TABLE_NAME	TABLE_TYPE
1	Adventure Works2014	Sales	vStoreWithContacts	VIEW
2	Adventure Works2014	Production	ScrapReason	BASE TABLE
3	Adventure Works2014	Sales	vStoreWithAddresses	VIEW
4	Adventure Works2014	Purchasing	vVendorWithContacts	VIEW
5	Adventure Works2014	HumanResources	Shift	BASE TABLE
6	Adventure Works2014	Purchasing	vVendorWithAddresses	VIEW
7	Adventure Works2014	Production	ProductCategory	BASE TABLE
8	Adventure Works2014	Purchasing	ShipMethod	BASE TABLE
9	Adventure Works2014	Production	ProductCostHistory	BASE TABLE
10	Adventure Works2014	Production	ProductDescription	BASE TABLE
11	Adventure Works2014	Sales	ShoppingCartItem	BASE TABLE
12	Adventure Works2014	Production	ProductDocument	BASE TABLE
13	Adventure Works2014	dbo	DatabaseLog	BASE TABLE
14	Adventure Works2014	Production	ProductInventory	BASE TABLE
15	Adventure Works2014	Sales	SpecialOffer	BASE TABLE
16	Adventure Works2014	dbo	ErrorLog	BASE TABLE
17	Adventure Works2014	Production	ProductListPriceHistory	BASE TABLE

```
SELECT * FROM INFORMATION_SCHEMA.TABLES;
```


More Examples

```
SELECT * FROM Sales.Customer;
```

	CustomerID	PersonID	StoreID	TerritoryID	AccountNumber	rowguid
1	1	NULL	934	1	AW00000001	3F5AE95E-B87D-4AED-95B4-C3797AFCEB74F
2	2	NULL	1028	1	AW00000002	E552F657-A9AF-4A7D-A645-C429D6E02491
3	3	NULL	642	4	AW00000003	130774B1-DB21-4EF3-98C8-C104BCD6ED6D
4	4	NULL	932	4	AW00000004	FF862851-1DAA-4044-BE7C-3E85583C054D
5	5	NULL	1026	4	AW00000005	83905BDC-6F5E-4F71-B162-C98DA069F38A
6	6	NULL	644	4	AW00000006	1A92DF88-BFA2-467D-BD54-FCB9E647FDD7
7	7	NULL	930	1	AW00000007	03E9273E-B193-448E-9823-FE0C44AEED78
8	8	NULL	1024	5	AW00000008	801368B1-4323-4BFA-8BEA-5B5B1E4BD4A0
9	9	NULL	620	5	AW00000009	B900BB7F-23C3-481D-80DA-C49A5BD6F772
10	10	NULL	928	6	AW00000010	CDB6698D-2FF1-4FBA-8F22-60AD1D11DABD

(19820 row(s) affected)

```
SELECT LastName FROM Person.Person;
```

	LastName
1	Abbas
2	Abel
3	Abercrombie
4	Abercrombie
5	Abercrombie
6	Abolrous
7	Abolrous
8	Acevedo
9	Achong
10	Ackerman
11	Ackerman
12	Adams
13	Adams
14	Adams

The FROM Clause

- The FROM clause is the very first query clause that is logically processed
- We specify the names of the tables that we want to query and table operators that operate on those tables

Note: If an identifier is irregular - for example, if it has embedded spaces or special characters, starts with a digit, or is a reserved Keyword - we have to delimit it

We can delimit identifiers in SQL Server in a couple of ways:

- double quotes - for example, "Order Details"
- square brackets - for example, [Order Details]

The WHERE Clause

- In the WHERE clause, we specify a predicate or logical expression to filter the rows returned by the FROM phase
- Only rows for which the logical expression evaluates to TRUE are returned by the WHERE phase to the subsequent logical query processing phase
- The WHERE clause has significance when it comes to query performance
 - Based on the filter expression, SQL Server evaluates the use of indexes to access the required data - get the required data with much less work compared to applying full table scans

The WHERE Clause (Cont.)

- The logical expression evaluates to **TRUE** are returned by the **WHERE** phase, actually, the query filters is “*accept TRUE*”
 - **WHERE** phase filters out both **FALSE** and **UNKNOWN**
 - Conversely, **CHECK** constraints is “*reject FALSE*”

- Example:

```
SELECT ProductID, Color
FROM   Production.Product
WHERE  Color <> 'Black'
```

	ProductID	Color
1	320	Silver
2	321	Silver
3	332	Silver
4	351	Silver
5	352	Silver
6	461	Silver
7	679	Silver
8	706	Red

```
SELECT ProductID, Color
FROM   Production.Product
WHERE  Color <> 'Black' OR Color IS NULL
```

	ProductID	Color
1	1	NULL
2	2	NULL
3	3	NULL
4	4	NULL
5	316	NULL
6	320	Silver
7	321	Silver
8	323	NULL

Three-valued Predicate Logic in T-SQL

- T-SQL uses three-valued predicate logic, where logical expressions can evaluate to **TRUE**, **FALSE**, or **UNKNOWN**.
- A logical expression involving only existing or present values evaluates to either **TRUE** or **FALSE**
- When the logical expression involves a missing value, it evaluates to **UNKNOWN**
- One of the tricky aspects of **UNKNOWN** is that when we negate it, we still get **UNKNOWN**
- An expression comparing two **NULL** marks (**NULL = NULL**) evaluates to **UNKNOWN**

The GROUP BY Clause

- The GROUP BY clause is used to aggregate information.
- For example, we want to know how many parts were ordered in a given set of orders

```
SELECT SalesOrderID, OrderQty
FROM Sales.SalesOrderDetail
WHERE SalesOrderID IN (43660, 43670, 43672);
```

- We can make use of the GROUP BY clause with an aggregator

```
SELECT SalesOrderID, SUM(OrderQty)
FROM Sales.SalesOrderDetail
WHERE SalesOrderID IN (43660, 43670, 43672)
GROUP BY SalesOrderID;
```

SalesOrderID	OrderQty
43660	1
43660	1
43670	1
43670	2
43670	2
43670	1
43672	6
43672	2
43672	1

(9 row(s) affected)

SalesOrderID	SUM(OrderQty)
43660	2
43670	6
43672	9

(3 row(s) affected)

The GROUP BY Clause (Cont.)

- When using a **GROUP BY** clause, all the columns in the **SELECT** list must either be aggregates (**SUM**, **MIN/MAX**, **AVG**, and so on) or columns included in the **GROUP BY** clause
- Likewise, if using an aggregate in the **SELECT** list, the **SELECT** list must only contain aggregates, or there must be a **GROUP BY** clause
- We can also group based on multiple columns. To do this we just add a comma and the next column name

The GROUP BY Clause (Cont.)

- All phases subsequent to the **GROUP BY** phase must operate on groups as opposed to operating on individual rows
- Each group is ultimately represented by a single row in the final result of the query
- All expressions processed in phases subsequent to the **GROUP BY** phase are required to guarantee returning a scalar per group

The GROUP BY Clause (Cont.)

- Expressions based on elements that participate in the **GROUP BY** list meet the requirement
 - By definition each *group* has only one *unique* occurrence of each GROUP BY element

- For example:

```
SELECT Color, YEAR(SellStartDate)
FROM Production.Product
WHERE Color <> 'Black'
GROUP BY Color, YEAR(SellStartDate)
```

Color	YEAR(...)
Blue	2011
Blue	2013
Grey	2012
Multi	2011
Multi	2012
Red	2008
Red	2011
Red	2012

Example: Aggregating Data

```
SELECT CustomerID, SalesPersonID, COUNT(*)
FROM Sales.SalesOrderHeader
WHERE CustomerID <= 11010
GROUP BY CustomerID, SalesPersonID
```

CustomerID	SalesPersonID	
11000	NULL	3
11001	NULL	3
11002	NULL	3
11003	NULL	3
11004	NULL	3
11005	NULL	3
11006	NULL	3
11007	NULL	3
11008	NULL	3
11009	NULL	3
11010	NULL	3

(11 row(s) affected)

Aggregates (Cont.)

- Elements that do not participate in the **GROUP BY** list are allowed only as inputs to an aggregate function
 - **AVG**: This one is for computing averages
 - **MIN/MAX**: These grab the minimum and maximum amounts for each grouping for a selected column
 - **COUNT**(Expression | *): The COUNT(*) function is about counting the rows in a query

Aggregates (Cont.)

- If an attribute that does not participate in the GROUP BY list and not as an input to an aggregate function, we get an error
 - No guarantee that the expression will return a single value per group
- Example:

```
SELECT    Color, YEAR(SellStartDate), ProductID
FROM      Production.Product
WHERE     Color <> 'Black'
GROUP BY  Color, YEAR(SellStartDate)
```

Column 'Production.Product.ProductID' is invalid in the select list because it is not contained in either an aggregate function or the GROUP BY clause.

Aggregates (Cont.)

- Note that all aggregate functions *ignore NULL* marks with one exception - COUNT(*)
- For example, consider a group of five rows called QTY
 - *COUNT(*) would return 5*
 - *COUNT(QTY) would return 4*
- **DISTINCT** keyword to handle only distinct occurrences of known values
 - *COUNT(DISTINCT QTY) would return 2*
 - *SUM(QTY) would return 60, and SUM(DISTINCT QTY) would return 40*

QTY
30
10
NULL
10
10

Groups with the HAVING Clause

- The **HAVING** clause is used only when there is also a **GROUP BY** in the query
 - Whereas the **WHERE** clause is applied to each row before it even has a chance to become part of a group
 - The **HAVING** clause is applied to the aggregated value for that group

WHERE: All of the conditions have been against specific *rows* before grouping.

HAVING: Place conditions on the *groups* themselves. In other words, apply the condition after the *groups* are fully accumulated.

The HAVING Clause

- With the **HAVING** clause, we can specify a predicate to filter groups as opposed to filtering individual rows
- Only groups for which the logical expression in the **HAVING** clause evaluates to **TRUE** are returned by the **HAVING** phase. (FALSE or UNKNOWN are filtered out)
- Because the HAVING clause is processed after the rows have been grouped, we can refer to *aggregate functions* in the logical expression. For example, HAVING COUNT(*) > 1

Example

```
CREATE TABLE HumanResources.Employee2 (  
    EmployeeID INT NOT NULL,  
    ManagerID INT NULL REFERENCES HumanResources.Employee2 (EmployeeID) ,  
    JobTitle NVARCHAR(50) NOT NULL,  
    LastName NVARCHAR(50) NOT NULL,  
    FirstName NVARCHAR(50) NOT NULL,  
    CONSTRAINT PK_EmployeeID PRIMARY KEY CLUSTERED ( EmployeeID ASC )  
);
```

INSERT INTO

```
HumanResources.Employee2 (EmployeeID, ManagerID, JobTitle, LastName,  
    FirstName) VALUES  
(1, NULL, 'CEO', 'Smith', 'Hunter'),  
(2, 1, 'CFO', 'Jones', 'Drew'),  
(3, 1, 'COO', 'Lenzy', 'Sheila'),  
(4, 1, 'CTO', 'Huntington', 'Karla'),  
(5, 4, 'VP of Engineering', 'Gutierrez', 'Ron'),  
(8, 5, 'VP of Engineering', 'Gutierrez', 'Ron'),  
(9, 5, 'Software Engineer', 'Bray', 'Marky'),  
(10, 5, 'Data Architect', 'Cheechov', 'Robert'),  
(11, 5, 'Software Engineer', 'Gale', 'Sue'),  
(6, 4, 'VP of Professional Services', 'Cross', 'Gary'),  
(7, 4, 'VP of Security', 'Lebowski', 'Jeff');
```


Example: The HAVING Clause

```
SELECT ManagerID, COUNT(*)
FROM   HumanResources.Employee2
GROUP BY ManagerID;
```

ManagerID	COUNT (*)
-----	-----
NULL	1
1	3
4	3
5	4

```
SELECT ManagerID, COUNT(*)
FROM   HumanResources.Employee2
WHERE  EmployeeID != 5
GROUP BY ManagerID;
```

ManagerID	COUNT (*)
-----	-----
NULL	1
1	3
4	2
5	4

```
SELECT ManagerID, COUNT(*)
FROM   HumanResources.Employee2
WHERE  EmployeeID != 5
GROUP BY ManagerID
HAVING COUNT(*) > 3;
```

ManagerID	COUNT (*)
-----	-----
5	4

The SELECT Clause

- The **SELECT** clause is where we specify the attributes that we want to return in the result table of the query
- T-SQL supports a couple of other forms with which we can alias expressions
 - <expression> **AS** <alias>,
 - AVG(UnitPrice) **AS** Price
 - <expression> <alias>,
 - AVG(UnitPrice) Price
 - <alias> = <expression>,
 - Price = AVG(UnitPrice)

Alias

- Column alias: representing derived and constant columns

```
SELECT    CategoryID, AVG(UnitPrice) Price
FROM      Products
GROUP BY  CategoryID
ORDER BY  Price;
```

Column alias can be used
in ORDER BY

```
SELECT ProductName, UnitPrice * 0.9 Discount
FROM    Products
WHERE   UnitPrice * 0.9 > 20;
```

Column alias can NOT be
used in WHERE or HAVING
clause (SQL Server)

- Table alias: commonly used in table joins and sub-queries

```
SELECT ProductName, CategoryName
FROM    Products AS p, Categories c
Where   p.CategoryID = c.CategoryID
```

"AS" is optional

If an alias is assigned, it must be
used instead of the original table
name

Alias Symbol

- Column alias: use [] or ' '

```
SELECT    ProductName, UnitPrice * 0.9 AS 'Discount Price'
FROM      Products
ORDER BY  'Discount Price';
```

```
SELECT    ProductName, UnitPrice * 0.9 AS [Discount Price]
FROM      Products
ORDER BY  [Discount Price];
```

- Table alias: use []

```
SELECT ProductName, CategoryName
FROM    Products AS [table p], Categories c
WHERE   [table p].CategoryID = c.CategoryID
```

The SELECT Clause (Cont.)

- The **SELECT** clause is processed after the **FROM**, **WHERE**, **GROUP BY**, and **HAVING** clauses
 - Aliases in the **SELECT** clause do not exist as far as clauses that are processed before the **SELECT** clause

```
SELECT Color, YEAR(SellStartDate) AS Year
FROM    Production.Product
WHERE    Year > 2007
```

```
Msg 207, Level 16, State 1, Line 3
Invalid column name 'Year'.
```

- Use of the same expression YEAR(...) in the query

```
SELECT Color, YEAR(SellStartDate) AS Year
FROM    Production.Product
WHERE    YEAR(SellStartDate) > 2007
```

The SELECT Clause - Result Set

- In the relational model, operations on relations are based on relational algebra and result in a relation (a set).
- In SQL, a **SELECT** query is not guaranteed to return a true set - namely, unique rows with no guaranteed order.
 - without a key, uniqueness of rows is not guaranteed, which case the table isn't a set; it's a multiset or a bag
- The term “result set” is often used to describe the output of a **SELECT** query
 - but a result set doesn't necessarily qualify as a set in the mathematical sense

The DISTINCT and ALL Predicates

- SQL provides the means to guarantee uniqueness in the result of a SELECT statement in the form of a DISTINCT clause that removes duplicate rows

```
SELECT DISTINCT Color, YEAR(SellStartDate) AS Year
FROM           Production.Product
WHERE          YEAR(SellStartDate) > 2007
```

Color	Year
-----	----
NULL	2008
NULL	2012
NULL	2013
...	

- ALL predicate says to include every row. ALL is the default for any SELECT statement.

```
SELECT ALL      Color, YEAR(SellStartDate) AS Year
FROM           Production.Product
WHERE          YEAR(SellStartDate) > 2007
```

Color	Year
-----	----
NULL	2008
NULL	2008
NULL	2008
...	

The SELECT Clause (Cont.)

- SQL supports the use of an asterisk (*) in the **SELECT** list to request all attributes from the queried tables
- Such use of an asterisk is a bad programming practice in most cases
 - Any schema changes applied to the table - such as adding or removing columns - might result in failures in the client application
- It is still *not allowed* to refer to a column alias that was created in the *same* SELECT clause

```
SELECT Color,  
       YEAR(SellStartDate) AS OrderYear,  
       OrderYear + 1      AS NextYear  
FROM   Production.Product
```


The ORDER BY Clause

- The **ORDER BY** clause allows us to sort the rows in the output for presentation purposes
- In terms of logical query processing, **ORDER BY** is the very last clause to be processed
- We can use any combination of columns in our **ORDER BY** clause
- The **DESC** (**ASC**) keyword tells SQL Server that the **ORDER BY** should work in descending (ascending) order

```

SELECT Name, SalesPersonID
FROM   Sales.Store
WHERE  Name BETWEEN 'g' AND 'j'
        AND SalesPersonID > 283
ORDER BY SalesPersonID, Name DESC;

```

Example: ORDER BY

Name	SalesPersonID
-----	-----
Inexpensive Parts Shop	286
Ideal Components	286
Helpful Sales and Repair Service	286
Helmets and Cycles	286
Global Sports Outlet	286
Gears and Parts Company	286
Irregulars Outlet	288
Hometown Riding Supplies	288
Good Bicycle Store	288
Global Bike Retailers	288
Instruments and Parts Company	289
Instant Cycle Store	290
Impervious Paint Company	290
Hiatus Bike Tours	290
Getaway Inn	290
(15 row(s) affected)	

The ORDER BY Clause (Cont.)

- A table has no guaranteed order, it is a set or multiset, a set has no order
- A table without specifying an **ORDER BY** clause, the query returns a table result, and SQL Server is free to return the rows in the output in any order
- The only way to guarantee that the rows in the result are sorted is to explicitly specify an **ORDER BY** clause
- If we do specify an **ORDER BY** clause, the result cannot qualify as a table - what standard SQL calls a *cursor* - a nonrelational result with order guaranteed among rows

The ORDER BY Clause (Cont.)

- T-SQL allows us to specify ordinal positions of columns in the ORDER BY clause, based on the order in which the columns appear in the SELECT list

```
SELECT Color, YEAR(SellStartDate) AS OrderYear,  
FROM Production.Product  
ORDER BY COLOR, YEAR(SellStartDate)
```

```
FROM Production ...  
ORDER BY 1, 2
```

- T-SQL allows us to specify elements in the ORDER BY clause that do not appear in the SELECT clause
- However, when DISTINCT is specified, we are restricted in the ORDER BY list only to elements that appear in the SELECT list

The ORDER BY Clause (Cont.)

- T-SQL allows us to specify elements in the ORDER BY clause that do not appear in the SELECT clause

```
SELECT    Color, YEAR(SellStartDate) AS OrderYear
FROM      Production.Product
ORDER BY  ProductID
```

- However, when **DISTINCT** is specified, we are restricted in the ORDER BY list only to elements that appear in the SELECT list.

```
SELECT DISTINCT Color, YEAR(SellStartDate) AS OrderYear
FROM            Production.Product
ORDER BY        ProductID
```

```
ORDER BY Color | ORDER BY OrderYear | ORDER BY Color, OrderYear
```

The TOP Filter

Color	Year
-----	----
NULL	2008
NULL	2008
NULL	2008
NULL	2008
NULL	2008

- The **TOP** option allows us to limit the number or percentage of rows that the query returns.

```
SELECT TOP (5) Color, YEAR(SellStartDate) AS OrderYear
FROM          Production.Product
ORDER BY      ProductID
```

- If **DISTINCT** is specified in the **SELECT** clause, the **TOP** filter is evaluated after duplicate rows have been removed
- The **PERCENT** keyword calculates the number of rows to return based on a percentage of the number of qualifying rows
- If we include **WITH TIES**, additional rows will be included if their values match or tie, the values of the loose row

The OFFSET-FETCH Filter

- Standard SQL defines a TOP-like filter called **OFFSET-FETCH** that does support skipping capabilities
- The **OFFSET-FETCH** filter is considered part of the **ORDER BY** clause, which normally serves a presentation ordering purpose
- By using the **OFFSET** clause, we can indicate how many rows to skip, and by using the **FETCH** clause, we can indicate how many rows to filter after the skipped rows

The OFFSET-FETCH Filter

- By using the **OFFSET** clause, we can indicate how many rows to skip, and by using the **FETCH** clause, we can indicate how many rows to filter after the skipped rows.

```
[ORDER BY <column list> [ASC|DESC]  
[OFFSET <count expression> {ROW|ROWS}  
  [FETCH {FIRST|NEXT} <count expression> {ROW|ROWS} ONLY]]
```

```
SELECT    Color, YEAR(SellStartDate) AS OrderYear  
FROM      Production.Product  
ORDER BY  ProductID  
OFFSET 10 ROWS FETCH NEXT 5 ROWS ONLY;
```

Color	Year
-----	----
Black	2008
NULL	2008
NULL	2008
NULL	2008
NULL	2008

Window Functions

- A window function is a function that, for each row in the underlying query, operates on a window of rows and computes a scalar result value
- The window of rows is defined by using an **OVER** clause
- Window functions are very profound and allow us to address a wide variety of needs
- The syntax of window function and **OVER** Clause:

```
WindowFunction (<expression>) OVER (  
    [ PARTITION BY <expression list> ]  
    [ ORDER BY <order list> ]  
)
```

Window Functions

- A window function operates on a set of rows exposed to it by a clause called **OVER**
- For each row in the underlying query, the **OVER** clause exposes to the function a subset of the rows from the underlying query's result set
- The **OVER** clause can restrict the rows in the window
 - by using the **PARTITION BY** subclause, and
 - by using the **ORDER BY** subclause to define ordering for the calculation

Window Functions

- The ROW_NUMBER window function assigns unique, sequential, incrementing integers to the rows in the result within the respective partition, based on the indicated ordering
- For example,

```
SELECT    Color, YEAR(SellStartDate) AS OrderYear, ListPrice,  
          ROW_NUMBER() OVER (PARTITION BY YEAR(SellStartDate)  
                              ORDER BY ListPrice) AS NUMBER  
FROM      Production.Product  
WHERE     ListPrice > 0  
ORDER BY  YEAR(SellStartDate)
```

- Partitions the window by YEAR(SellStartDate)
- Order the records in each window by ListPrice
- Assign the row number to each record in each window

Window Functions

Color	OrderYear	ListPrice	Rownum
-----	-----	-----	-----
NULL	2008	133.34	1
NULL	2008	133.34	2
NULL	2008	133.34	3
...			
NULL	2008	196.92	9
Black	2008	1431.50	10
Red	2008	1431.50	11
Multi	2011	8.99	1
White	2011	9.50	2
White	2011	9.50	3
Blue	2011	34.99	4
...			
Red	2011	3578.27	71
Red	2011	3578.27	72
NULL	2012	13.99	1
NULL	2012	19.99	2
...			
(304 row(s) affected)			

Summary

- To put it all together, the following list presents the logical order in which all clauses discussed so far are processed:
 1. FROM
 2. WHERE
 3. GROUP BY
 4. SELECT
 - Expressions
 - DISTINCT
 5. ORDER BY
 - TOP/OFFSET-FETCH

Exercises

1. Write a query against the Production.Product table that returns ProductID, Color and SellStartDate that SellStartDate in 2007.
2. Write a query against the Sales.SalesOrderDetail table that returns orders with total value (OrderQty* UnitPrice) greater than 10,000, sorted by total value.

	SalesOrderID	SalesOrderDetailID	CarrierTrackingNumber	OrderQty	ProductID	SpecialOfferID	UnitPrice	UnitPriceDiscount	LineTotal
1	43659	1	4911-403C-98	1	776	1	2024.994	0.00	2024.994000
2	43659	2	4911-403C-98	3	777	1	2024.994	0.00	6074.982000
3	43659	3	4911-403C-98	1	778	1	2024.994	0.00	2024.994000
4	43659	4	4911-403C-98	1	771	1	2039.994	0.00	2039.994000
5	43659	5	4911-403C-98	1	772	1	2039.994	0.00	2039.994000
6	43659	6	4911-403C-98	2	773	1	2039.994	0.00	4079.988000
7	43659	7	4911-403C-98	1	774	1	2039.994	0.00	2039.994000
8	43659	8	4911-403C-98	3	714	1	2		
9	43659	9	4911-403C-98	1	716	1	2		
10	43659	10	4911-403C-98	6	709	1	5.70	0.00	34.200000
11	43659	11	4911-403C-98	2	712	1	5.1865	0.00	10.373000
12	43659	12	4911-403C-98	4	711	1	20.1865	0.00	80.746000

SalesOrderDetail

Exercises

3. Write a query against the Sales.SalesOrderHeader table that returns the three TerritoryID with the highest average freight in 2012 (based on OrderDate).

	SalesOrderID	OrderDate	SalesOrderNumber	AccountNumber	CustomerID	SalesPersonID	TerritoryID	SubTotal	Freight
1	43659	2011-05-31 00:00:00.000	SO43659	10-4020-000676	29825	279	5	20565.6206	616.0984
2	43660	2011-05-31 00:00:00.000	SO43660	10-4020-000117	29672	279	5	1294.2529	38.8276
3	43661	2011-05-31 00:00:00.000	SO43661	10-4020-000442	29734	282	6	32726.4786	985.553
4	43662	2011-05-31 00:00:00.000	SO43662	10-4020-000227	29994	282	6	28832.5289	867.2389
5	43663	2011-05-31 00:00:00.000	SO43663	10-4020-000510	29565	276	4	419.4589	12.5838
6	43664	2011-05-31 00:00:00.000	SO43664	10-4020-000397	29898	280	1	24432.6088	732.81
7	43665	2011-05-31 00:00:00.000	SO43665	10-4020-000146	29580	283	1	14352.7713	429.9821
8	43666	2011-05-31 00:00:00.000	SO43666	10-4020-000511	30052	276	4	5056.4896	151.9921
9	43667	2011-05-31 00:00:00.000	SO43667	10-4020-000646	29974	277	3	6107.082	183.1626
10	43668	2011-05-31 00:00:00.000	SO43668	10-4020-000514	29614	282	6	35944.1562	1081.8...
11	43669	2011-05-31 00:00:00.000	SO43669	10-4020-000578	29747	283	1	714.7043	22.0367
12	43670	2011-05-31 00:00:00.000	SO43670	10-4020-000504	29566	275	3	6122.082	183.6126
13	43671	2011-05-31 00:00:00.000	SO43671	10-4020-000200	29890	283	1	8128.7876	244.0042
14	43672	2011-05-31 00:00:00.000	SO43672	SalesOrderHeader				6124.182	183.6257
15	43673	2011-05-31 00:00:00.000	SO43673					3746.2015	111.8629
16	43674	2011-05-31 00:00:00.000	SO43674					2624.382	78.7315

ProductID	int	Not null
Name	Name (user-defined type) nvarchar(50)	Not null
ProductNumber	nvarchar(25)	Not null
MakeFlag	Flag (user-defined type) bit	Not null
FinishedGoodsFlag	Flag (user-defined type) bit	Not null
Color	nvarchar(15)	null
SafetyStockLevel	smallint	Not null
ReorderPoint	smallint	Not null
StandardCost	money	Not null
ListPrice	money	Not null
Size	nvarchar(5)	null
SizeUnitMeasureCode	nchar(3)	null
WeightUnitMeasureCode	nchar(3)	null
Weight	decimal (8,2)	null
DaysToManufacture	int	Not null
ProductLine	nchar(2)	null
Class	nchar(2)	null
Style	nchar(2)	null
ProductSubcategoryID	smallint	null
ProductModelID	int	null
SellStartDate	datetime	Not null
SellEndDate	datetime	null
DiscontinuedDate	datetime	null
rowguid	uniqueidentifier ROWGUIDCOL	Not null
ModifiedDate	datetime	Not null

SalesOrderID	int	Not null	Primary key. Foreign key to SalesOrderHeader.SalesOrderID.
SalesOrderDetailID	int	Not null	Primary key. A sequential number used to ensure data uniqueness..
CarrierTrackingNumber	nvarchar(25)	Null	Shipment tracking number supplied by the shipper.
OrderQty	smallint	Not null	Quantity ordered per product.
ProductID	int	Not null	Product sold to customer. Foreign key to Product.ProductID.
SpecialOfferID	int	Not null	Promotional code. Foreign key to SpecialOffer.SpecialOfferID.
UnitPrice	money	Not null	Selling price of a single product.
UnitPriceDiscount	money	Not null	Discount amount.
LineTotal	Computed as OrderQty * UnitPrice	Not null	Per product subtotal.
rowguid	uniqueidentifier ROWGUIDCOL	Not null	ROWGUIDCOL number uniquely identifying the row. Used to support a merge replication sample.
ModifiedDate	datetime	Not null	Date and time the row was last updated.

SalesOrderID	int	Not null	Sales.SalesOrderHeader
RevisionNumber	tinyint	Not null	
OrderDate	datetime	Not null	
DueDate	datetime	Not null	
ShipDate	datetime	Null	
Status	tinyint	Not null	
OnlineOrderFlag	Flag (user-defined type) bit	Not null	
SalesOrderNumber	nvarchar(25)	Not null	
PurchaseOrderNumber	OrderNumber (user-defined type) nvarchar(25)	Null	
AccountNumber	AccountNumber (user-defined type) nvarchar(15)	Null	
CustomerID	int	Not null	
ContactID	int	Not null	
SalesPersonID	int	Null	
TerritoryID	int	Null	
BillToAddressID	int	Not null	
ShipToAddressID	int	Not null	
ShipMethodID	int	Not null	
CreditCardID	int	Null	
CreditCardApprovalCode	varchar(15)	Null	
CurrencyRateID	int	Null	
SubTotal	money	Not null	
TaxAmt	money	Not null	
Freight	money	Not null	
TotalDue	Computed as SubTotal + TaxAmt + Freight	Not null	
Comment	nvarchar(128)	Null	
rowguid	uniqueidentifier ROWGUIDCOL	Not null	50
ModifiedDate	datetime	Not null	