

CS108: Advanced Database

- Database Programming

Lecture 06:
Subqueries

Overview

- Subqueries
- Table Expression
 - Derived Tables
 - Common Table Expression
 - View
 - Inline Table-Valued Function

Subquery

- SQL supports writing queries within queries, or nesting queries.
- The outermost query is a query whose result set is returned to the caller and is known as the outer query.
- The inner query is a query whose result is used by the outer query and is known as a subquery.
- A subquery can be either *self-contained* or *correlated*.
- A self-contained subquery has no dependency on the outer query that it belongs to, whereas a correlated subquery does.
- A subquery can be single-valued, multivalued, or table-valued.

Subquery Search Conditions

- A subquery usually appears as part of a search condition in the **WHERE** or **HAVING** clause.
- The type of subquery search conditions are:
 - *Subquery comparison test* - Compares the value of an expression with a single value produced by a subquery.
 - *Subquery set membership test* - Checks whether the value of an expression matches one of the set of values produced by a subquery.
 - *Existence test* - Tests whether a subquery produces any rows of query results.
 - *Quantified comparison test* - Compares the value of an expression with each of the sets of values produced by a subquery.

Self-Contained Subqueries

- Every subquery has an outer query that it belongs to.
- Self-contained subqueries are subqueries that are independent of the outer query that they belong to.

```
SELECT <SELECT list>
FROM   <SomeTable>
WHERE  <SomeColumn> {=, <>, <, <=, >, >=} (
        SELECT <single column>
        FROM   <SomeTable>
        WHERE  <condition that results in only one row returned>)
```

```
SELECT <SELECT list>
FROM   <SomeTable>
WHERE  <SomeColumn> IN (
        SELECT <single column>
        FROM   <SomeTable>
        [WHERE <condition>])
```

Self-Contained Subqueries (cont.)

- Every subquery has an outer query that it belongs to.
- Self-contained subqueries are subqueries that are independent of the outer query that they belong to.

```
SELECT <SELECT list>
FROM   <SomeTable>
WHERE  <SomeColumn> EXIST (
        SELECT <single column>
        FROM   <SomeTable>
        WHERE  <condition>)
```

```
SELECT <SELECT list>
FROM   <SomeTable>
WHERE  <SomeColumn> {ALL, ANY, SOME} {=, <>, <, <=, >, >=} (
        SELECT <single column>
        FROM   <SomeTable>
        [WHERE <condition>])
```

The Subquery Comparison Test

- A *scalar* subquery is a subquery that returns a single value - regardless of whether it is self-contained.
- Such a subquery can appear anywhere in the outer query where a *single-valued* expression can appear (such as WHERE or SELECT).
- For example: Suppose that we need to query the Orders table in the AdventureWorks database and return information about the order that has the maximum order ID in the table.

- We could accomplish the task by using a variable

```
DECLARE @maxid AS INT = (SELECT MAX(SalesOrderID)
                           FROM Sales.SalesOrderHeader);

SELECT SalesOrderID, OrderDate, SalesPersonID, CustomerID
FROM Sales.SalesOrderHeader
WHERE SalesOrderID = @maxid;
```

SalesOrderID	OrderDate	SalesPersonID	CustomerID
75123	2014-06-30 00:00:00.000	NULL	18759

- We can substitute the technique that uses a variable with an embedded subquery.

```
SELECT SalesOrderID, OrderDate, SalesPersonID, CustomerID
FROM Sales.SalesOrderHeader
WHERE SalesOrderID = (SELECT MAX(SalesOrderID)
                       FROM Sales.SalesOrderHeader);
```


Self-Contained Scalar Subquery

- For a scalar subquery to be valid, it **MUST** return no more than one value.
- If a scalar subquery can return more than one value, it might fail at run time.

```
SELECT SalesOrderID
FROM Sales.SalesOrderHeader
WHERE SalesPersonID = (SELECT P.BusinessEntityID
                        FROM Sales.SalesPerson AS P
                        WHERE P.TerritoryID = 5);
```

```
...
71950
71951
```

```
(429 row(s) affected)
```

```
SELECT SalesOrderID
FROM Sales.SalesOrderHeader
WHERE SalesPersonID = (SELECT P.BusinessEntityID
                        FROM Sales.SalesPerson AS P
                        WHERE P.TerritoryID = 1);
```

Subquery returned more than 1 value. This is not permitted when the subquery follows =, !=, <, <=, >, >= or when the subquery is used as an expression.

The Subquery Comparison Test

- If a scalar subquery returns no value, it returns a NULL.
- Recall that a comparison with a NULL yields UNKNOWN and that query filters do not return a row for which the filter expression evaluates to UNKNOWN.
- For example, the Sales.SalesPerson table currently has no person whose TerritoryID = 50; therefore, the following query returns an empty set.

```
SELECT SalesOrderID
FROM Sales.SalesOrderHeader
WHERE SalesPersonID = (SELECT P.BusinessEntityID
                        FROM Sales.SalesPerson AS P
                        WHERE P.TerritoryID = 50);
```

(0 row(s) affected)

The Set Membership Test (IN)

- A multivalued subquery is a subquery that returns multiple values as a single column.
- The subquery set membership test (**IN**) compares a single data value with a column of data values produced by a subquery and returns a **TRUE** result if the data value matches one of the values in the column.
- Note that the use of **DISTINCT** in the subquery isn't strictly necessary. If the same record appears multiple times in the subquery results instead of only once, the outer query yields the same results.

- For example:

```
SELECT SalesOrderID
FROM Sales.SalesOrderHeader
WHERE SalesPersonID IN (SELECT P.BusinessEntityID
                        FROM Sales.SalesPerson AS P
                        WHERE P.TerritoryID = 1);
```

...
71943

(424 row(s) affected)

- Similarly, we can solve this problem using joins.

```
SELECT SalesOrderID
FROM Sales.SalesOrderHeader AS S
JOIN Sales.SalesPerson      AS P
  ON S.SalesPersonID = P.BusinessEntityID
WHERE P.TerritoryID = 1
```

...
71943

(424 row(s) affected))

- Many querying problems that we can solve with either subqueries or joins.

Subqueries and Joins

- Many queries write using subqueries can also write as multitable queries, or joins.
- Conversely, many queries with subqueries cannot be translated into an equivalent join.
- For example, list the ID which SubTotal have above average SubTotal.

```
SELECT SalesOrderID
FROM Sales.SalesOrderHeader
WHERE SubTotal > (SELECT AVG(SubTotal) FROM Sales.SalesOrderHeader)
```

- In this case, the inner query is a summary query and the outer query is not, so there is no way the two queries can be combined into a single join.

The Set Membership Test (IN)

- As with any other predicate, we can negate the **IN** predicate with the **NOT** logical operator.
- For example, the following query returns customers who did not place any orders.

```
SELECT CustomerID
FROM Sales.Customer
WHERE CustomerID NOT IN
      (SELECT CustomerID FROM Sales.SalesOrderHeader)
```

...
700
701

(701 row(s) affected)

- Note that best practice is to qualify the subquery to exclude **NULL** marks.
- If the subquery includes **NULL** marks then the finally result will be an empty set.

- Note that best practice is to qualify the subquery to exclude NULL marks.

```
SELECT CustomerID
FROM Sales.Customer
WHERE CustomerID NOT IN (
    SELECT CustomerID FROM Sales.SalesOrderHeader
    UNION ALL
    SELECT NULL )
```

CustomerID

(0 row(s) affected)

- Remember that T-SQL uses three-valued logic. The WHERE only accepts true.
- $X \text{ NOT IN } (n_1, n_2, \dots, \text{NULL}, \dots, n_N)$
- $= \text{NOT } (X \text{ IN } (n_1, n_2, \dots, \text{NULL}, \dots, n_N))$
- $= \text{NOT } (X = n_1 \text{ OR } X = n_2 \text{ OR } X = \text{NULL} \text{ OR } \dots \text{ OR } X = n_N)$
- $= \text{NOT } (X = n_1 \text{ OR } X = n_2 \text{ OR } \text{UNKNOWN} \text{ OR } \dots \text{ OR } X = n_N)$
- $= (X \neq n_1 \text{ AND } X \neq n_2 \dots \text{ AND } X \neq n_N) \text{ AND } \text{NOT UNKNOWN}$
- $\in \{ \text{FALSE}, \text{UNKNOWN} \}$ (row always be filtered out)

Example: Set Membership Test (IN)

- Example: Using a nested **SELECT** to find missing records

```
SELECT    O.Description
FROM      Sales.SpecialOffer O
LEFT JOIN Sales.SpecialOfferProduct P
          ON P.SpecialOfferID = O.SpecialOfferID
WHERE     P.SpecialOfferID IS NULL;
```

Description

Volume Discount over 60
(1 row(s) affected)

```
SELECT    Description
FROM      Sales.SpecialOffer O
WHERE     O.SpecialOfferID NOT IN
          (SELECT SpecialOfferID FROM Sales.SpecialOfferProduct);
```

Description

Volume Discount over 60
(1 row(s) affected)

Example: Set Membership Test (IN)

- Example: Return all individual CustomerIDs that are missing between the minimum and maximum in the table.
- The Nums table is useful in this example.
 - The Nums table contains a sequence of integers, starting with 1, with no gaps.
 - To return all missing IDs from the Customer table, query the Nums table filter only
 - The numbers that are between the minimum and maximum in the Customer table and
 - The numbers do not appear in the set of CustomerIDs in the Customer table.

Example: Set Membership Test (IN)

- Example: Return all individual CustomerIDs that are missing between the minimum and maximum in the table.

```
SELECT N
FROM (SELECT O1.N + O2.N * 10 + O3.N * 100 + O4.N*1000 + O5.N*10000 AS N
      FROM      (VALUES (0) , (1) , (2) , (3) , (4) , (5) , (6) , (7) , (8) , (9)) AS O1 (N)
      CROSS JOIN (VALUES (0) , (1) , (2) , (3) , (4) , (5) , (6) , (7) , (8) , (9)) AS O2 (N)
      CROSS JOIN (VALUES (0) , (1) , (2) , (3) , (4) , (5) , (6) , (7) , (8) , (9)) AS O3 (N)
      CROSS JOIN (VALUES (0) , (1) , (2) , (3) , (4) , (5) , (6) , (7) , (8) , (9)) AS O4 (N)
      CROSS JOIN (VALUES (0) , (1) , (2) , (3) , (4) , (5) , (6) , (7) , (8) , (9)) AS O5 (N)
      ) AS Num
WHERE N BETWEEN (SELECT MIN(CustomerID) FROM Sales.Customer AS O)
              AND (SELECT MAX(CustomerID) FROM Sales.Customer AS O)
              AND N NOT IN (SELECT      CustomerID FROM Sales.Customer AS O);
```

N

...

10989

10999

(10298 row(s) affected)

Correlated Subqueries

- Correlated subqueries are subqueries that refer to attributes from the table that appears in the outer query.
- This means that the subquery is dependent on the outer query and cannot be invoked independently. (difficult to debug)

- For example, we want to know the first order of each customer

```
-- Get a list of customers and the date of their first order
SELECT      CustomerID, MIN(OrderDate) AS OrderDate
INTO        #MinOrderDates
FROM        Sales.SalesOrderHeader
GROUP BY    CustomerID;

-- Do something additional with that information
SELECT      A.CustomerID, A.SalesOrderID, A.OrderDate
FROM        Sales.SalesOrderHeader AS A
JOIN        #MinOrderDates AS T
  ON        A.CustomerID = T.CustomerID
  AND       A.OrderDate  = T.OrderDate
ORDER BY    A.CustomerID;

DROP TABLE #MinOrderDates;
```

```
...
30118          47378          2012-07-31 00:00:00.000
(19134 row(s) affected)
```

- If we want this to run in a single query, we must find a way to look up each individual customer.
- We can do this by making use of an *inner* query that performs a lookup based on the current CustomerID in the *outer* query.
- We then must return a value back out to the outer query so it can match things up based on the earliest order date.

```
SELECT  OUT.CustomerID, OUT.SalesOrderID, OUT.OrderDate
FROM    Sales.SalesOrderHeader AS OUT
WHERE   OUT.OrderDate = (
        SELECT MIN(IN.OrderDate)
        FROM    Sales.SalesOrderHeader AS IN
        WHERE   IN.CustomerID = OUT.CustomerID )
ORDER BY OUT.CustomerID;
```

Example: Correlated Subqueries

- Example: suppose that we need to query the Sales.SalesOrderHeader view and return for each order the percentage that the current order value is of the total values of all of the customer's orders.

```
SELECT    O1.SalesOrderID, O1.CustomerID,
          CAST ( 100. * O1.SubTotal / (SELECT SUM(O2.SubTotal)
                                     FROM    Sales.SalesOrderHeader AS O2
                                     WHERE   O2.CustomerID = O1.CustomerID )
              AS NUMERIC(5, 2) ) AS [%]
FROM      Sales.SalesOrderHeader AS O1
ORDER BY  CustomerID, SalesOrderID
```

```
...
58928      30118      12.84
65221      30118      8.71
71803      30118      11.74
```

(31465 row(s) affected)

The Existence Test (EXISTS)

- SQL supports a predicate called **EXISTS** that accepts a subquery as input and returns **TRUE** if the subquery returns any rows and **FALSE** otherwise.
- Notice that the **EXISTS** search condition doesn't really use the results of the subquery at all.
- It merely tests to see whether the subquery produces any results.
 - For this reason, SQL relaxes the rule that “subqueries must return a single column of data” and allows us to use the **SELECT** * form in the subquery of an **EXISTS** test.

- Example: find customers who did place any orders.
- Using EXISTS

```
SELECT C.CustomerID, C.AccountNumber
FROM Sales.Customer AS C
WHERE EXISTS (
    SELECT *
    FROM Sales.SalesOrderHeader AS O
    WHERE O.CustomerID = C.CustomerID )
```

- Using IN

```
SELECT C.CustomerID, C.AccountNumber
FROM Sales.Customer AS C
WHERE C.CustomerID IN (SELECT O.CustomerID
    FROM Sales.SalesOrderHeader AS O
    WHERE O.CustomerID = C.CustomerID )
```

- Using JOIN

```
SELECT DISTINCT C.CustomerID, C.AccountNumber
FROM Sales.Customer AS C
LEFT OUTER JOIN Sales.SalesOrderHeader AS O
    ON O.CustomerID = C.CustomerID
WHERE O.SalesOrderID IS NOT NULL
```


The Existence Test (EXISTS)

- Another aspect of the **EXISTS** predicate that is interesting to note is that unlike most predicates in SQL, **EXISTS** uses two-valued logic and not three-valued logic.
 - **TRUE** - If the subquery returns any rows
 - **FALSE** - otherwise.

```
SELECT 1
WHERE EXISTS (SELECT NULL)           -- (1 row(s) affected)

SELECT 1
WHERE EXISTS (SELECT 1)              -- (1 row(s) affected)

SELECT 1
WHERE EXISTS (SELECT TOP (0) 1)      -- (0 row(s) affected)

SELECT 1
WHERE EXISTS (SELECT 1 WHERE -1 = 1) -- (0 row(s) affected)
```

Quantified Tests (ANY and ALL)

- SQL provides two quantified tests, **ANY** and **ALL**, that extend comparison operators, such as greater than.
- Both of these tests compare a data value with the column of data values produced by a subquery.
- For example: List the salespeople who have taken an order that represents more than 10 percent of their quota.

```
SELECT *  
FROM Sales.SalesPerson AS S  
WHERE (.1 * S.SalesQuota) < ANY (  
    SELECT TotalDue  
    FROM Sales.SalesOrderHeader AS O  
    WHERE O.SalesPersonID = S.BusinessEntityID )
```

- Conceptually, the main query tests each row one by one.

The ANY/SOME Test

- The ANY/SOME test is used in conjunction with one of the six SQL comparison operators (=, <>, <, <=, >, >=) to compare a single test value with a column of data values produced by a subquery.
- To perform the test, SQL uses the specified comparison operator to compare the test value with each data value in the column, one at a time.
- If **any** of the individual comparisons yields a **TRUE** result, the **ANY/SOME** test returns a **TRUE** result.

The ANY/SOME Test (cont.)

- The ANSI/ISO SQL standard specifies these detailed rules describing the results of the ANY/SOME test when the test value is compared with the column of subquery results:
 - If the subquery produces an empty column of query results, the ANY/SOME test returns FALSE - no value is produced by the subquery for which the comparison test holds.
 - If the comparison test is TRUE for at least one of the data values in the column, then the ANY/SOME search condition returns TRUE - indeed some value is produced by the subquery for which the comparison test holds.

The ANY/SOME Test (cont.)

- The ANSI/ISO SQL standard specifies these detailed rules describing the results of the ANY/SOME test when the test value is compared with the column of subquery results:
 - If the comparison test is FALSE for every data value in the column, then the ANY/SOME search condition returns FALSE - no value which the comparison test holds.
 - If the comparison test is not TRUE for any data value in the column, but it is NULL (unknown) for one or more of the data values, then the ANY/SOME search condition returns NULL.

The ANY/SOME Test (cont.)

```
SELECT 'NULL'      --NULL
WHERE NOT EXISTS (
    SELECT 1
    WHERE NOT (1 = ANY ( SELECT 2 UNION SELECT NULL))
           OR  (1 = ANY ( SELECT 2 UNION SELECT NULL)))
```

```
SELECT 'NOT NULL'  --NOT NULL
WHERE EXISTS (
    SELECT 1
    WHERE NOT (1 = ANY ( SELECT 2 UNION SELECT 100))
           OR  (1 = ANY ( SELECT 2 UNION SELECT 100)))
```

- If the comparison test is **not TRUE** for **any** data value in the column, but it is **NULL** (unknown) for one or more of the data values, then the **ANY/SOME** search condition returns **NULL**.

The ANY/SOME Test (cont.)

- The **ANY** comparison operator can be very tricky to use in practice, especially in conjunction with the inequality (<>) comparison operator.
- Here is an example that shows the problem: List the ID of all the SalesPerson who do not have an order.
- It's tempting to express this query as shown in this example:

```
SELECT S.BusinessEntityID
FROM Sales.SalesPerson AS S
WHERE S.BusinessEntityID <> ANY (SELECT O.SalesPersonID
                                     FROM Sales.SalesOrderHeader AS O)
```


The ALL Test

- Like the **ANY** test, the **ALL** test is used in conjunction with one of the six SQL comparison operators (**=**, **<>**, **<**, **<=**, **>**, **>=**) to compare a single test value with a column of data values produced by a subquery.
- To perform the test, SQL uses the specified comparison operator to compare the test value with each data value in the column, one at a time.
- If **all** of the individual comparisons yield a **TRUE** result, the **ALL** test returns a **TRUE** result.

The ALL Test (cont.)

- The ANSI/ISO SQL standard specifies rules describing the results of the **ALL** test:
 - If the subquery produces an **empty** column of query results, the **ALL** test returns **TRUE** - the comparison test does hold for every value produced by the subquery.
 - If the comparison test is **TRUE for every** data value in the column, then the **ALL** search condition returns **TRUE**.

The ALL Test (cont.)

- The ANSI/ISO SQL standard specifies rules describing the results of the ALL test:
 - If the comparison test is FALSE for any data value in the column, then the ALL search condition returns FALSE.
 - If the comparison test is not FALSE for any data value in the column, but it is NULL for one or more of the data values, then the ALL search condition returns NULL.

Subqueries in the HAVING Clause

- Subqueries can also be used in the **HAVING** clause of a query.
- When a subquery appears in the **HAVING** clause, it works as part of the row group selection performed by the **HAVING** clause.
- Because the subquery is evaluated once for each row group, all outer references in the correlated subquery must be single-valued for each row group.
 - Either be a reference to a grouping column of the outer query
 - Or be contained within a column (aggregate) function.

Subqueries in the HAVING Clause

- Example: List the salespeople whose average TotalDue for selling product No. 910 is at least as big as that salesperson's overall average TotalDue.

```
SELECT    H1.SalesPersonID
FROM      Sales.SalesOrderHeader AS H1
JOIN      Sales.SalesOrderDetail AS D
  ON      H1.SalesOrderID = D.SalesOrderID
WHERE     D.ProductID = 910
GROUP BY  H1.SalesPersonID
HAVING    AVG(H1.TotalDue) >= (
                                SELECT AVG(H2.TotalDue)
                                FROM      Sales.SalesOrderHeader AS H2
                                WHERE     H2.SalesPersonID = H1.SalesPersonID )
```

Subqueries Example

- Example: Returning Previous or Next Values

```
SELECT SalesOrderID, OrderDate, SalesPersonID, CustomerID,
       (SELECT MAX(SalesOrderID)
        FROM Sales.SalesOrderHeader AS O2
        WHERE O2.SalesOrderID < O1.SalesOrderID) AS PrevSalesOrderID
FROM Sales.SalesOrderHeader AS O1
```

SalesOrderID	OrderDate	SalesPersonID	CustomerID	PrevSalesOrderID
-----	-----	-----	-----	-----
43659	2011-05-31 00:00:00.000	279	29825	NULL
43660	2011-05-31 00:00:00.000	279	29672	43659
...				

```
SELECT SalesOrderID, OrderDate, SalesPersonID, CustomerID,
       (SELECT MIN(SalesOrderID)
        FROM Sales.SalesOrderHeader AS O2
        WHERE O2.SalesOrderID > O1.SalesOrderID) AS NextSalesOrderID
FROM Sales.SalesOrderHeader AS O1
```

...				
75122	2014-06-30 00:00:00.000	NULL	15868	75123
75123	2014-06-30 00:00:00.000	NULL	18759	NULL

■ Example: Running Aggregates

```
SELECT  YEAR(OrderDate) AS OrderYear, SUM(TotalDue) AS Qty
INTO    #OrderYear
FROM    Sales.SalesOrderHeader
WHERE   SalesPersonID = 282
GROUP BY YEAR(OrderDate)

SELECT  OrderYear, Qty,
        (SELECT SUM(O2.Qty)
         FROM    #OrderYear AS O2
         WHERE   O2.OrderYear <= O1.OrderYear) AS RunQty
FROM    #OrderYear AS O1
ORDER BY OrderYear

DROP TABLE #OrderYear
```

OrderYear	Qty	RunQty
2011	1323328.6346	1323328.6346
2012	2070323.5014	3393652.136
2013	2112546.1213	5506198.2573
2014	1177338.401	6683536.6583

The background features five circles arranged in two rows. The top row contains three circles: the leftmost is an outline, and the two on the right are solid light purple. The bottom row contains three circles: the two on the left are solid light purple, and the rightmost is an outline.

Table Expression

Table Expressions

- A table expression is a named query expression that represents a valid relational table.
- We can use table expressions in data manipulation statements much like we use other tables.
- Microsoft SQL Server supports four types of table expressions:
 - Derived tables
 - Common table expressions (CTEs)
 - Views
 - Inline table-valued functions (inline TVFs)

Derived Tables

- Derived tables (also known as table subqueries) are defined in the **FROM** clause of an outer query.
- Their scope of existence is the outer query.
- As soon as the outer query is finished, the derived table is gone.
- We specify the query that defines the derived table within parentheses, followed by the **AS** clause and the derived table name.

Derived Tables

- We specify the query that defines the derived table within parentheses, followed by the **AS** clause and the derived table name.
- Example:

```
SELECT *  
FROM    (SELECT ProductID, Name, ListPrice  
         FROM    Production.Product  
         WHERE   Color = 'Black') AS BlackColorProduct
```

	ProductID	Name	ListPrice
1	317	LL Crankarm	0.00
2	318	ML Crankarm	0.00
3	319	HL Crankarm	0.00
4	322	Chainring	0.00
5	680	HL Road Frame - Black, 58	1431.50
6	708	Sport-100 Helmet, Black	34.99
7	722	LL Road Frame - Black, 58	337.22
8	723	LL Road Frame - Black, 60	337.22
9	724	LL Road Frame - Black, 62	337.22

Derived Tables

- A query must meet three requirements to be valid to define a table expression of any kind:
 - *Order is not guaranteed.*
 - A table expression is supposed to represent a relational table, and the rows in a relational table have no guaranteed order.
 - For this reason, standard SQL disallows an **ORDER BY** clause in queries that are used to define table expressions, unless the **ORDER BY** serves another purpose besides presentation (**OFFSET-FETCH** and **TOP** filter).

Derived Tables

- A query must meet three requirements to be valid to define a table expression of any kind:
 - *All columns must have names.*
 - All columns in a table must have names; therefore, we must assign column aliases to all expressions in the **SELECT** list of the query that is used to define a table expression.
 - *All column names must be unique.*
 - All column names in a table must be unique; therefore, a table expression that has multiple columns with the same name is invalid.

Using Arguments

- In the query that defines a derived table, we can refer to arguments.
- The arguments can be local variables and input parameters to a routine such as a stored procedure or function.

```
DECLARE @Color AS NVARCHAR(100) = 'Black';

SELECT *
FROM    (SELECT ProductID, Name, ListPrice
        FROM    Production.Product
        WHERE   Color = @Color) AS BlackColorProduct;
```

Nesting

- If we need to define a derived table by using a query that itself refers to a derived table, we end up nesting derived tables.
- Nesting of derived tables is a result of the fact that a derived table is defined in the **FROM** clause of the outer query and not separately.
- Example:

OrderYear	NumCusts
-----	-----
2013	11095
2014	10354
2012	3162

```
SELECT OrderYear, NumCusts
FROM    (SELECT OrderYear, COUNT(DISTINCT CustomerID) AS NumCusts
        FROM    (SELECT YEAR(OrderDate) AS OrderYear, CustomerID
                FROM    Sales.SalesOrderHeader) AS D1
        GROUP BY OrderYear) AS D2
WHERE   NumCusts > 2000
```

Multiple References

- Another problematic aspect of derived tables stems from the fact that derived tables are defined in the **FROM** clause of the outer query and not prior to the outer query.
- As far as the **FROM** clause of the outer query is concerned, the derived table doesn't exist yet; therefore, if we need to refer to multiple instances of the derived table, we can't.
- Instead, we have to define multiple derived tables based on the same query.

Multiple References

```
SELECT Cur.OrderYear,  
       Cur.NumCusts AS CurNumCusts, Prv.NumCusts AS PrvNumCusts,  
       Cur.NumCusts - Prv.NumCusts AS Growth  
FROM   (SELECT OrderYear, COUNT(DISTINCT CustomerID) AS NumCusts  
        FROM   (SELECT YEAR(OrderDate) AS OrderYear, CustomerID  
                 FROM   Sales.SalesOrderHeader) AS D1  
        GROUP BY OrderYear) AS Cur  
LEFT JOIN (SELECT OrderYear, COUNT(DISTINCT CustomerID) AS NumCusts  
          FROM   (SELECT YEAR(OrderDate) AS OrderYear, CustomerID  
                   FROM   Sales.SalesOrderHeader) AS D1  
          GROUP BY OrderYear) AS Prv  
ON Cur.OrderYear = Prv.OrderYear + 1  
ORDER BY OrderYear
```

OrderYear	CurNumCusts	PrvNumCusts	Growth
-----	-----	-----	-----
2011	1406	NULL	NULL
2012	3162	1406	1756
2013	11095	3162	7933
2014	10354	11095	-741

Common Table Expressions

- Common table expressions (CTEs) are another standard form of table expression very similar to derived tables, yet with a couple of important advantages.
- CTEs are defined by using a **WITH** statement and have the following general form.

```
WITH <CTE_Name>[ (<target_column_list>)]  
AS  
(  
    <inner_query_defining_CTE>  
)  
<outer_query_against_CTE>;
```

- The inner query defining the CTE must follow all requirements mentioned earlier to be valid to define a table expression.

Assigning Column Aliases in CTEs

- CTEs also support two forms of column aliasing - inline and external.
- For the inline form, specify <expression> **AS** <column_alias>.
- For the external form, specify the target column list in parentheses immediately after the CTE name.

```
WITH <CTE_Name>
AS
(
    SELECT [<expression> AS <column_alias>] ...
    ...
)
<outer_query_against_CTE>;
```

Example: Find the TotalCost (TotalDue + Freight) for customers

```
WITH Orders AS (  
    SELECT SalesOrderID, CustomerID, TotalDue + Freight AS TotalCost  
    FROM    Sales.SalesOrderHeader  
)  
SELECT C.CustomerID, Orders.SalesOrderID, Orders.Total  
FROM    Sales.Customer AS C  
JOIN    Orders ON C.CustomerID = Orders.CustomerID;
```

■ Or

CustomerID	SalesOrderID	TotalCost
-----	-----	-----
29825	43659	23769.3323
29672	43660	1496.1564
29734	43661	37851.3542
...		

```
WITH Orders ([Order ID], [Customer ID], Total)  
AS (  
    SELECT SalesOrderID, CustomerID, TotalDue + Freight AS TotalCost  
    FROM    Sales.SalesOrderHeader  
)  
SELECT C.CustomerID, Orders.[Order ID], Orders.Total  
FROM    Sales.Customer AS C  
JOIN    Orders ON C.CustomerID = Orders.[Customer ID];
```

Views and Inline TVFs

- Derived tables and CTEs have a very limited scope, which is the single-statement scope. Derived tables and CTEs are *not reusable*.
- Views and inline table-valued functions (inline TVFs) are two *reusable* types of table expressions; their definitions are stored as database objects.
- In most other respects, views and inline TVFs are treated like derived tables and CTEs.

Views

- A view is a reusable table expression whose definition is stored in the database.
- Note that the general recommendation to avoid using **SELECT *** has specific relevance in the context of views.
 - The columns are enumerated in the compiled form of the view, and new table columns will **NOT** be automatically added to the view.

```
CREATE VIEW <view name> [(<column name list>)]  
[WITH [ENCRYPTION] [[,] SCHEMABINDING] [[,] VIEW_METADATA]]  
AS  
    <SELECT statement>  
[WITH CHECK OPTION] [;]
```

Views

- A view is a reusable table expression whose definition is stored in the database.
- Example:

```
IF OBJECT_ID('Production.BLACKColorProduct') IS NOT NULL
    DROP VIEW Production.BLACKColorProduct;
GO

CREATE VIEW Production.BLACKColorProduct
AS
    SELECT ProductID, Name, ListPrice
    FROM Production.Product
    WHERE Color = 'Black';
GO

SELECT *
FROM Production.BLACKColorProduct
```

Views and the ORDER BY Clause

- The query that we use to define a view *must meet* all requirements mentioned earlier with respect to table expressions in the context of derived tables.
- The view should *not guarantee* any order to the rows, all view columns must have names, and all column names must be unique.
- Remember that a presentation **ORDER BY** clause is *not* allowed in the query defining a table expression because there's no order among the rows of a relational table.

- An attempt to create an ordered view is absurd because it violates fundamental properties of a relation as defined by the relational model.

```
ALTER VIEW Production.BLACKColorProduct
AS
    SELECT    ProductID, Name, ListPrice
    FROM      Production.Product
    WHERE     Color = 'Black'
    ORDER BY  Name
GO
```

```
Msg 1033, Level 15, State 1, Procedure BLACKColorProduct, Line 73
The ORDER BY clause is invalid in views, inline functions, derived
tables, subqueries, and common table expressions, unless TOP, OFFSET
or FOR XML is also specified.
```

- The error message indicates that SQL Server allows the **ORDER BY** clause in three exceptional cases - when the **TOP**, **OFFSET-FETCH** option is used.

Views and the ORDER BY Clause

- Because T-SQL allows an **ORDER BY** clause in a view when **TOP** or **OFFSET-FETCH** is also specified, but we can not create “ordered views” in SQL Server.
- One of the ways to try to achieve this is by using **TOP** (100) **PERCENT**, like the following.

```
ALTER VIEW Production.BLACKColorProduct
AS
    SELECT TOP(100) PERCENT ProductID, Name, ListPrice
    FROM    Production.Product
    WHERE   Color = 'Black'
    ORDER BY Name
GO
```

- Even though the code is technically valid and the view is created, we should be aware that because the query is used to define a table expression, the **ORDER BY** clause here is only guaranteed to serve the logical filtering purpose for the **TOP** option.
- If we query the view and don't specify an **ORDER BY** clause in the outer query, presentation order is **not** guaranteed.

```
SELECT Name, ListPrice, ProductID
FROM    Production.BLACKColorProduct
```

Name	ListPrice	ProductID
-----	-----	-----
LL Crankarm	0.00	317
ML Crankarm	0.00	318
HL Crankarm	0.00	319
Chainring	0.00	322
HL Road Frame - Black, 58	1431.50	680
...		

View Option

- The **ENCRYPTION** option indicates that SQL Server will internally store the text with the definition of the object in an obfuscated format.
- The **SCHEMABINDING** option is available to views and UDFs; it binds the schema of referenced objects and columns to the schema of the referencing object. It indicates that referenced objects (base table or tables) cannot be dropped and that referenced columns cannot be dropped or altered.
- The purpose of **CHECK OPTION** is to prevent modifications through the view that conflict with the view's filter - assuming that one exists in the query defining the view.

Inline Table-Valued Functions

- Inline TVFs are reusable table expressions that support input parameters. In all respects except for the support for input parameters, inline TVFs are similar to views.
- Inline TVF is the user defined function that return a table.

```
CREATE FUNCTION <function name>
( [ <@parameter name> [AS] <data type>
  [ = <default value> [READONLY]]
  [ ,...n ] ] )
RETURNS {<scalar type> | TABLE [( <table definition> )]}
[ WITH [ENCRYPTION] | [SCHEMABINDING] |
  [RETURNS NULL ON NULL INPUT | CALLED ON NULL INPUT] |
  [EXECUTE AS { CALLER|SELF|OWNER|<'user name'> } ]
] [AS] { EXTERNAL NAME <external method> |
BEGIN
  [<function statements>]
  {RETURN <type as defined in RETURNS clause> | RETURN (<SELECT statement>)}
END } [;]
```

- For example, the following code creates an inline TVF called GetCustOrders in the database.

```
IF OBJECT_ID('GetCustOrders') IS NOT NULL
DROP FUNCTION GetCustOrders;
GO
```

```
CREATE FUNCTION GetCustOrders(@cid AS INT)
RETURNS TABLE
AS
```

```
    RETURN
```

```
        SELECT SalesOrderID, CustomerID,
               SalesPersonID, OrderDate, ShipMethodID,
               ShipDate, TerritoryID, Freight, SubTotal
```

```
        FROM    Sales.SalesOrderHeader
```

```
        WHERE   CustomerID = @cid;
```

```
GO
```

```
SELECT O.SalesOrderID, O.CustomerID
FROM    GetCustOrders(30085) AS O;
```

SalesOrderID	CustomerID
-----	-----
46937	30085
47968	30085
49044	30085
50196	30085
51796	30085
...	

- As with other tables, we can refer to an inline TVF as part of a join.
- For example, the following query joins the inline TVF returning customer 30085's orders with the SalesOrderDetails table, matching customer 30085's orders with the related order lines.

```
SELECT O.SalesOrderID, O.CustomerID, OD.ProductID, OD.OrderQty
FROM   GetCustOrders(30085)      AS O
JOIN   Sales.SalesOrderDetail    AS OD
ON     O.SalesOrderID = OD.SalesOrderID;
```

SalesOrderID	CustomerID	ProductID	OrderQty
-----	-----	-----	-----
46937	30085	809	1
47968	30085	817	1
47968	30085	815	1
47968	30085	859	4
49044	30085	814	1
...			
(22 row(s) affected)			

The APPLY Operator

- The **APPLY** operator is a very powerful table operator.
- Like all table operators, this operator is used in the **FROM** clause of a query.
- The two supported types of **APPLY** operator are **CROSS APPLY** and **OUTER APPLY** - similar to cross join, outer join.
- **CROSS APPLY** implements only one logical query processing phase, whereas **OUTER APPLY** implements two (add a **NULL** marks to a resulting empty set).

The APPLY Operator

- The **APPLY** operator operates on two input tables, the second of which can be a table expression; we'll refer to them as the “left” and “right” tables.
- The right table is usually a derived table or an inline TVF.
- The **CROSS APPLY** operator implements one logical query processing phase - it applies the right table expression to each row from the left table and produces a result table with the unified result sets.

The CROSS APPLY Operator

- The **CROSS APPLY** operator is very similar to a cross join.

```
SELECT N
FROM (SELECT O1.N + O2.N * 10 AS N
      FROM (VALUES (0), (1), (2), (3), (4), (5), (6), (7), (8), (9)) AS O1 (N)
      CROSS JOIN (VALUES (0), (1), (2), (3), (4), (5), (6), (7), (8), (9)) AS O2 (N)
      ) AS Num;
```

N

0
...
99

```
SELECT N
FROM (SELECT O1.N + O2.N * 10 AS N
      FROM (VALUES (0), (1), (2), (3), (4), (5), (6), (7), (8), (9)) AS O1 (N)
      CROSS APPLY (VALUES (0), (1), (2), (3), (4), (5), (6), (7), (8), (9)) AS O2 (N)
      ) AS Num;
```

- However, with the **CROSS APPLY** operator, the right table expression can represent a different set of rows per each row from the left table, unlike in a join.

- The **CROSS APPLY** is intended to enable joining to TVFs.

```
SELECT N
FROM (SELECT O1.N + O2.N * 10 AS N
      FROM (VALUES (0), (1), (2), (3), (4), (5), (6), (7), (8), (9)) AS O1(N)
      CROSS JOIN (VALUES (O1.N + 0), (O1.N + 1), (O1.N + 2)) AS O2(N)
      ) AS Num
```

```
SELECT N
FROM (SELECT O1.N + O2.N * 10 AS N
      FROM (VALUES (0), (1), (2), (3), (4), (5), (6), (7), (8), (9)) AS O1(N)
      CROSS APPLY (VALUES (O1.N + 0), (O1.N + 1), (O1.N + 2)) AS O2(N)
      ) AS Num
```

- We can use a derived table in the right side, and in the derived table query refer to attributes from the left side.
- Alternatively, when we use an inline TVF, we can pass attributes from the left side as input arguments.

- The **CROSS APPLY** is intended to enable joining to TVFs.

```
CREATE FUNCTION CreateNUM(@N AS INT)
RETURNS TABLE
AS
    RETURN
        SELECT N
        FROM    (VALUES (@N+0) , (@N+1) , (@N+2) , (@N+3) , ... , (@N+9)) AS O(N)
GO
```

```
SELECT O1.N + O2.N * 10
FROM    CreateNUM(0) AS O1
CROSS JOIN CreateNUM(0) AS O2
```

```
SELECT O1.N + O2.N * 10
FROM    CreateNUM(0) AS O1
CROSS JOIN CreateNUM(O1.N) AS O2
```

```
SELECT O1.N + O2.N * 10
FROM    CreateNUM(0) AS O1
CROSS APPLY CreateNUM(O1.N) AS O2
```

- The **CROSS APPLY** is intended to enable joining to TVFs.

```

SELECT SOH.CustomerID, SOH.OrderDate, SOH.TotalDue, CRT.RunningTotal
FROM   Sales.SalesOrderHeader AS SOH
CROSS APPLY (
    SELECT SUM(TotalDue) AS RunningTotal
    FROM   Sales.SalesOrderHeader AS RT
    WHERE  RT.CustomerID    = SOH.CustomerID
          AND RT.SalesOrderID <= SOH.SalesOrderID
) AS CRT
ORDER BY SOH.CustomerID, SOH.SalesOrderID

```

CustomerID	OrderDate	TotalDue	RunningTotal
-----	-----	-----	-----
11000	2011-06-21 00:00:00.000	3756.989	3756.989
11000	2013-06-20 00:00:00.000	2587.8769	6344.8659
11000	2013-10-03 00:00:00.000	2770.2682	9115.1341
11001	2011-06-17 00:00:00.000	3729.364	3729.364
11001	2013-06-18 00:00:00.000	2674.0227	6403.3867
11001	2014-05-12 00:00:00.000	650.8008	7054.1875
11002	2011-06-09 00:00:00.000	3756.989	3756.989
11002	2013-06-02 00:00:00.000	2535.964	6292.953
...			

The OUTER APPLY Operator

- The **OUTER APPLY** operator adds a second logical phase that identifies rows from the left side for which the right table expression returns an empty set, and it adds those rows to the result table as outer rows with **NULL** marks in the right side's attributes as placeholders.
- In a sense, this phase is similar to the phase that adds outer rows in a left outer join.

The OUTER APPLY Operator

- Use **OUTER APPLY** like a **LEFT OUTER JOIN** to return a row from the left even if there is nothing returned from the right.

```
SELECT  PRD.ProductID, S.SalesOrderID
FROM    Production.Product AS PRD
OUTER APPLY (
    SELECT TOP(2) SalesOrderID
    FROM      Sales.SalesOrderDetail AS SOD
    WHERE     SOD.ProductID = PRD.ProductID
    ORDER BY SalesOrderID) AS S
ORDER BY PRD.ProductID
```

ProductID	SalesOrderID
1	NULL
2	NULL
3	NULL
4	NULL
316	NULL
317	NULL
...	